

# JUST ENOUGH C FOR 243

*C is quirky, flawed, and an enormous success.*

— Dennis Ritchie, co-designer of C

*A C program is like a fast dance on a newly waxed dance floor by people carrying razors.*

— Waldi Ravens, NetBSD developer

This guide summarizes the C programming language as used in *EEE243 Applied Programming* at the Royal Military College of Canada. It's more of a quick reference than a tutorial, it deliberately leaves things out, and it occasionally over-simplifies, but it's enough to get you started.

The guide occasionally refers to Joseph H. Silverman's *C Reference Card*, from [www.math.brown.edu/~jhs/](http://www.math.brown.edu/~jhs/). The *Card* provides more breadth but less explanation. For further information, see the course text or the Wikipedia articles on C and its syntax.

In this guide C code is written like `this`. Anything written *like this* in a code sample would need to be replaced with real C code in a program.

---

## Contents

<b>Contents</b>	<b>1</b>	comparison.....	5
<b>C program structure</b>	<b>2</b>	Boolean operators (&&,   , !) .....	5
main program file .....	2	array index expression.....	5
modules or libraries.....	2	pointer operators (address and dereference).....	6
header files.....	2	pointer arithmetic.....	6
module implementation files.....	2	function call expression.....	6
<b>Preprocessor directives</b>	<b>2</b>	ternary conditional expression ( a ? b : c ).....	6
#include.....	2	<b>Statements</b>	<b>6</b>
#define.....	2	expression statements .....	6
#ifdef and #ifndef.....	2	block statement (sequence) .....	7
<b>Basic types</b>	<b>2</b>	if and if-else statements (selection).....	7
integer types (including bool and char) .....	2	switch statement (selection) .....	7
unsigned integers.....	2	while statement (loop).....	7
floating point types.....	2	do-while statement (loop) .....	7
void.....	3	for statement (loop).....	7
<b>Arrays</b>	<b>3</b>	break.....	8
<b>Pointers</b>	<b>3</b>	continue .....	8
NULL.....	3	<b>Functions</b>	<b>8</b>
<b>Values and literal representations</b>	<b>3</b>	function prototype (declaration) .....	8
numbers .....	3	function definition .....	8
truth values .....	3	return statement .....	8
characters.....	3	static variables .....	8
escape sequence character literals.....	3	<b>Comments</b>	<b>9</b>
array literals.....	3	<b>Memory management</b>	<b>9</b>
character strings .....	3	sizeof .....	9
<b>Identifiers (names)</b>	<b>3</b>	malloc.....	9
<b>Variable and parameter declarations</b>	<b>3</b>	calloc .....	9
const.....	4	free .....	9
static.....	4	<b>printf and sprintf</b>	<b>9</b>
<b>enum definition and use</b>	<b>4</b>	literal characters .....	9
<b>struct definition and use</b>	<b>4</b>	escape sequences.....	9
<b>typedef definition and use</b>	<b>4</b>	value placeholders (conversion specifiers).....	9
<b>Expressions</b>	<b>5</b>	<b>Copyright and use</b>	<b>10</b>
operators and precedence.....	5		
assignment (=).....	5		
compound assignment (+, *= ...) .....	5		
modulus.....	5		
integer division.....	5		
increment and decrement (++ , --).....	5		



---

## C program structure

A C program will have a *main program file* and will almost always include one or more *modules* or *libraries*.

### main program file

Exactly one file in a program (whose name ends in .c) must have a function called `main`, which is the initial point of execution. The normal structure of a C program's main file is:

*includes of header files required by this program*  
*other preprocessor directives*

*global declarations, including prototypes for*  
*functions defined later in this file*

```
int main(void) {  
    local declarations  
    statements  
    return EXIT_SUCCESS;  
}
```

*function definitions*

### modules or libraries

There are two kinds of modules: those you and your team write as part of your program, and those provided to you. Some modules are part of the standard C distribution; others come from vendors or other third parties. A C module consists of a *header file*, whose name ends in .h, and an *implementation file* whose name ends in .c. The first part of the name must match. The implementation file may be pre-compiled and distributed in object code form.

#### header files

A header file contains declarations of functions and constants that are intended to be reused across multiple programs. It is always distributed in source code form. A header file is brought into your program using the `#include` directive.

#### module implementation files

Module implementation files provide the definitions of all functions declared in their corresponding header files. Every module is compiled separately; the resulting *object files* are combined by the *linker* to produce your program. Third party modules are often provided in pre-compiled form.

---

## Preprocessor directives

Preprocessor directives run prior to compilation and modify the text of the program itself.

### #include

The `#include` directive requires a file name as an argument. This will always be a C header file. The preprocessor inserts the contents of the header file at the point where the `#include` occurs. There are two forms:

```
#include <filename.h>  
    look for a file called filename.h in the system-  
    defined library directories  
#include "filename.h"  
    look for a file called filename.h in the same  
    directory as the current file
```

### #define

The `#define` directive is used for *replacement text*, typically *constants* (it can also define *replacement macros*, which we will not use in the course). For example

```
#define MAX_HEADROOM 2.3
```

will cause every occurrence of the characters `MAX_HEADROOM` in the program text to be replaced with the characters `2.3`

### #ifdef and #ifndef

The `#ifdef` and `#ifndef` directives allow for conditional compilation depending on whether a symbol has been defined using `#define`. A typical use of `#ifndef` is in header files, to prevent a header's contents from being included multiple times (which could result in multiple declarations of the same function). This is how the pattern would be used in a header file named `schroedinger.h`

```
#ifndef SCHROEDINGER_H_  
#define SCHROEDINGER_H_  
    code to be included exactly once  
#endif
```

If the symbol `SCHROEDINGER_H_` is not defined, then define it and include everything up to `#endif`. If this header is included a second time, `SCHROEDINGER_H_` will be defined, so the body of the `#ifndef` will be skipped. This assumes the symbol `SCHROEDINGER_H_` is globally unique in the program.

---

## Basic types

The C programming model is very close to raw hardware. Values are stored as bit sequences. Types define the number of bits a value occupies in memory and how those bits should be interpreted.

In C, the bit size (and therefore range and precision) of the primitive types depends on both the underlying hardware and the compiler implementation. Srsly.

### integer types (including bool and char)

<code>int</code>	signed integer, usually 16 or 32 bits — the default type for an integer literal
<code>long</code>	signed integer, usually longer than an <code>int</code>
<code>char</code>	a single ASCII character, normally 8 bits; yes, <code>char</code> is a kind of integer
<code>bool</code>	a truth value, either 1 (meaning true) or 0 (meaning false). <code>bool</code> , <code>true</code> and <code>false</code> are not actually part of the C language and are defined in the library <code>stdbool.h</code>

### unsigned integers

A signed integer of  $n$  bits can represent numbers from  $-2^{n-1}$  to  $2^{n-1}-1$ . So if an `int` is 16 bits long, a signed `int` can represent values from -32768 to 32767. If negative values are not required, the `unsigned` keyword can be added to the declaration to give us a larger positive range. A 16 bit `unsigned int` permits us to store values from 0 to 65535.

### floating point types

<code>float</code>	signed floating point number, normally IEEE 754 single-precision (32 bits)
<code>double</code>	signed floating point number, normally IEEE 754 double-precision (64 bits) — the default type for a floating point literal

## void

The C type `void` has two meanings. When used as the return type of a function, `void` means that the function does not return a value. A void pointer is a pointer to a value of undefined type, so a void pointer is a “universal pointer.”

---

## Arrays

A C *array* is a fixed-size region of memory that stores a sequence of values of some base type. The base type may be either a primitive type or a *derived type* (array, pointer, enum or struct). All elements must have the same type. Array types are indicated by square brackets `[]`. See **literal arrays**, **variable and parameter declarations**, and **array index expression**.

---

## Pointers

A *pointer* represents the location in memory at which a value is stored. Pointers are themselves values: a pointer is a special integer type whose size is equal to the hardware address length.

For examples of how to declare a pointer type, see **Variable and parameter declarations**. For examples of how to use pointers, see **pointer operators** and **pointer arithmetic**.

## NULL

NULL is a void pointer to memory location zero. This should not be confused with the null character `'\0'` (see **escape sequence character literals**). Both have value zero, but `'\0'` is a char (one byte wide) and NULL is a pointer (typically eight bytes wide).

---

## Values and literal representations

### numbers

42	integer (exactly 42) of type <code>int</code>
42L	integer of type <code>long</code>
42LU	integer of type <code>unsigned long</code>
42.7	floating point number (approximately 42.7) of type <code>double</code>
42.7F	floating point number of type <code>float</code>
4.275e2	exponential notation: equivalent to 427.5

### truth values

Include the library `stdbool.h` in your program to use the literal values `true` and `false` and the type `bool`.

In C anything that evaluates to zero is considered false, and anything non-zero is considered true.

### characters

<code>'c'</code>	single quotes denote characters; only characters in the ASCII range may be used
------------------	---

## escape sequence character literals

The following “escape sequences” may be used in character or string literals. There are others; these are the ones you will need most frequently.

<code>\n</code>	new line (in output, advances to the next line)
<code>\"</code>	double quote <code>"</code>
<code>\'</code>	single quote <code>'</code>
<code>\\</code>	backslash <code>\</code>
<code>\0</code>	the null character, used to terminate strings; ASCII value 0. Not to be confused with NULL
<code>%%</code>	a literal <code>%</code> character, needed only in <code>printf</code> format strings — elsewhere, a simple <code>%</code> is fine

## array literals

Array literals are written with the array values enclosed in curly braces and separated by commas and may be used *only* in array declaration expressions. Array literals may be nested.

<code>{ 1, 42, -3 }</code>	a three-element array of <code>int</code>
<code>{ { 1, 42, -3 }, { 12, 99, 14 } }</code>	a two-element array, each element of which is a three element <code>int</code> array
<code>{ 1, 42, -3, 12, 99, 14 }</code>	because C stores arrays in row-major order, this is exactly equivalent to the previous expression

## character strings

C uses null-terminated arrays of characters to represent strings. Literal strings are written with double quotes.

<code>"Hello"</code>	equivalent to the character array <code>{ 'H', 'e', 'l', 'l', 'o', '\0' }</code> (note the terminating null character)
----------------------	--

---

## Identifiers (names)

Identifiers in C are used to name variables, functions, function parameters, and types declared by `typedef`, `struct`, and `enum`.

A name must begin with a letter or an underscore, which may be followed by any combination of letters, digits, and underscores. Uppercase and lowercase are different, so `foo` and `F00` and `Foo` are three different names. Names may be any length; however, only the first 31 characters are guaranteed to be significant.

C forbids use of these *reserved words* for names: `auto`, `break`, `case`, `char`, `const`, `continue`, `default`, `do`, `double`, `else`, `entry`, `enum`, `extern`, `float`, `for`, `goto`, `if`, `int`, `long`, `register`, `return`, `short`, `signed`, `sizeof`, `static`, `struct`, `switch`, `typedef`, `union`, `unsigned`, `void`, `volatile`, and `while`.

---

## Variable and parameter declarations

All variables and function parameters (see **Function declarations**) in C programs must be *declared* before use. A declaration reserves space in memory for the variable or parameter, sufficient to hold a value of the declared type. Variable declarations may include initializations (see **assignment**) but parameter declarations may not. Here are some examples:

`int x`      the variable called `x` is of type `int`  
`int x = 5`   `x` is of type `int` and is initialized to 5  
`int a[4]`    `a` is an array of four `int` values. See the section on **array and string expressions**.  
`int a[]`     `a` is an array of `int`, length unknown. If this is a variable declaration, `a` must also be initialized here in order to determine its size; if `a` is a function parameter, this is equivalent to `int *a`  
`int *p`      `p` is a pointer to a value of type `int`. See the sections on **Pointers, pointer operations, and pointer arithmetic**  
`int *p[10]`   `p` is an array of 10 pointers to `int` values  
`int (*p)[10]`   `p` is a pointer to an array of 10 `int` values (this is different from the previous declaration)  
`char *s`     `s` is a pointer to a value of type `char`. This is the type of C string values.  
Bork `b`      `b` is of type `Bork`. Since `Bork` is not a standard C type, it must have been defined previously in the program via a `typedef`; see the section on **typedef**.

## const

The prefix `const` may be used to declare that a variable or function parameter is a constant value. If a declared constant is a variable, it must be initialized at the point of declaration.

```
const int x = 37
```

`x` is an integer constant with the value 37. Any attempt to modify the value of `x` will be flagged as illegal by the compiler.

## static

See the section on **static variables** in the **Functions** section.

---

## enum definition and use

The `enum` keyword can be used to define an *enumerated type*, a type that has a fixed set of named values. The syntax is

```
enum tag { values };
```

where *tag* is an identifier and *values* is a comma-separated list of identifiers. For example:

```
enum AccountKind { SAVINGS, CHEQUING, CREDIT };
// defines a new type, enum AccountKind, with the
// three legal values
```

To declare a variable of an enumerated type we use the full type name, in this case `enum Account`. For example:

```
enum AccountKind new_account = CREDIT;
// declares a variable new_account to be of type
// enum AccountKind and gives it an initial value
// of CREDIT
```

Each of the values in an enumerated type is treated as an integer, starting from zero. So, in our example, `CREDIT` is actually the integer 2. Always use the value names when working with `enum` types, never the integer equivalents — otherwise you defeat the purpose of having them.

---

## struct definition and use

Structures allow us to group related values into a named unit. The syntax is

```
struct tag { fields };
```

where *tag* is an identifier and the *field-list* is a semicolon-separated list of fields, each of which takes the form of a variable declaration, and may be of any defined type, including `struct` or enumerated types. For example:

```
struct Client {
    char name[30];
    int client_number;
    enum AccountKind account;
};
```

defines a new type, `struct Client`, with three fields called `name`, `client_number`, and `account`. Note the trailing semicolon.

To declare a variable of a `struct` type we use the full type name, in this case `struct Client`. For example

```
struct Client bob;
```

declares a variable `bob` to be of type `struct Client`

We can assign all values of a `struct` at once, but *only* on initialization.

```
struct Client bob = {"Bob Smith", 25, SAVINGS};
```

Anywhere else, we must access each field separately, using *dot notation*, like this

```
bob.account = CHEQUING;
```

---

## typedef definition and use

The `typedef` keyword can be used to give an alternate name to a type. The syntax is

```
typedef type new-name ;
```

The *type* can be any C type and the *new-name* can be any valid C identifier. For example:

```
typedef int Pounds;
```

defines the new type name `Pounds`, which has the same meaning as `int` and can be used anywhere `int` can be used

```
typedef struct Client BankClient;
```

defines a new type name `BankClient` with the same meaning as `struct Client`

Types defined via `typedef` can be used anywhere a type is expected.

It is common to combine `typedef` with `enum` or `struct` definition. The syntax is the same, with the `enum` or `struct` definition appearing as the *type*. In this case, the *tag* is optional. For example

```
typedef struct ClientTag {
    char name[30];
    int client_number;
    enum AccountKind account;
} Client ;
```

## Expressions

In C, an expression is anything that has a value (possibly void) when evaluated. Expressions may include variables, literal values, operators, and function calls. Expressions are *evaluated left to right*, according to *operator precedence rules*, and ultimately result in a single value.

### operators and precedence

For a list of operators and their precedence, see the *C Reference Card*. Normal arithmetic operators have their usual precedence, so the expression `6+3*-7` evaluates to `-15` and is equivalent to `6+(3*(-7))`. However, C operator precedence is *tricky*. If you're in doubt, add parentheses (...) to enforce order, but try not to overdo it. Explanations of some of the less common and less obvious operators are provided below.

### assignment (=)

Modifies the value of a variable.

`var = expression`

Evaluates *expression* and replaces the contents of *var* with the result. *var* must be a declared variable or an array index expression and it must be possible to coerce the type of *expression* to match the declared type of *var*. The value of an assignment is the value of its right-hand side.

### compound assignment (+=, \*= ...)

Compound assignment is a short form for the frequent case of updating a variable by performing some operation on its current value. C provides compound assignment forms for the operators `+`, `-`, `*`, `/`, and `%` as well as the bitwise operators.

`x *= 3 + 5` is equivalent to `x = x * (3 + 5)`

### modulus

The modulus operator `%` applies to integers only and evaluates to the remainder after integer division.

`27 % 7`      `6`, the remainder after dividing 27 by 7

### integer division

Dividing two integers gives an integer result, which is truncated, not rounded. So `2/3` evaluates to `0`. Floating point division works as expected and can be forced by making either the dividend or the divisor floating point, so `2.0/3` is approximately `0.6666667`.

### increment and decrement (++ , --)

C provides operators to increment (`++`) and decrement (`--`) variables by one. If the increment/decrement appears to the left of the variable in an expression, the operation is performed before the expression is evaluated. If it appears to the right, the operation is performed afterwards. If `x`, `y`, and `z` are integer variables, the code in the two columns below is equivalent:

<code>y = x++;</code>	<code>y = x;</code>
	<code>x = x + 1;</code>
<code>z = ++x;</code>	<code>x = x + 1;</code>
	<code>z = x;</code>

## comparison

The C comparison operators evaluate to `1` if the condition is true and `0` if the condition is false. Be careful not to mistakenly use `=` (assignment) when you mean `==` (equality comparison).

<code>==</code>	is equal to ( <code>a == b</code> is true if <code>a</code> is less than <code>b</code> )
<code>!=</code>	is not equal to
<code>&lt;</code>	is less than
<code>&gt;</code>	is greater than
<code>&lt;=</code>	is less than or equal to
<code>&gt;=</code>	is greater than or equal to

Since the comparison operators have precedence after the arithmetic operators, `5 < 4 + 3` is true, since the `4 + 3` is evaluated first.

Comparisons operate on two values at a time. So, perhaps surprisingly, `5 > 4 > 3` evaluates to `0` (false) since `5 > 4` evaluates to `1` and `1 > 3` is false. If you need to compare multiple values you must use the Boolean operators.

### Boolean operators (&&, ||, !)

The Boolean operators support complex truth value expressions. Watch out for the `&` and `|` bitwise operators! Using `&` when you mean `&&` is a frequent cause of error.

<code>!</code>	logical NOT or “negation.” <code>!a</code> is true if <code>a</code> is not true; <code>!a</code> is false if <code>a</code> is not false
<code>&amp;&amp;</code>	logical AND. <code>a &amp;&amp; b</code> is true if <code>a</code> is true and <code>b</code> is also true
<code>  </code>	logical OR. <code>a    b</code> is true if <code>a</code> is true or if <code>b</code> is true, including the case where both are true

The correct way to write the (legal but misleading) expression `5 > 4 > 3` is `5 > 4 && 4 > 3` which means “5 is greater than 4 AND 4 is greater than three” – which evaluates to `1` (true).

The `&&` and `||` operators are both “short circuiting,” which means that they stop evaluating as soon as the truth value of the whole expression can be known. So in `a && b`, if `a` is false then `b` will never be evaluated. Similarly, in `a || b`, if `a` is true then `b` will never be evaluated. This is important for expressions that cause side effects when evaluated.

### array index expression

The individual elements of arrays (including strings) can be accessed by *indexing*, which is written `expression[index]` where *expression* must evaluate to a pointer. (An array name is a constant pointer to the first value of the array.)

The index of the first element in an array is `0`. If the index refers to a position that doesn't exist in the matrix, C does not report an error but happily returns the value stored at the indicated memory location.

In the following examples, assume the initialization

```
int v[3] = { 4, 5, 6 };
char *s = "abcde";
int m[2][3] = { { 7, 8, 9 }, { 10, 11, 12 } };
```

`v[1]` is the second element of `v`, which is 5.  
`v[4]` whatever happens to be located in memory four elements past the start of `v`, which is generally garbage since `v` has no fifth element.  
[imdb.com/title/tt0119116/](http://imdb.com/title/tt0119116/)  
`s[1]` the second element of `s`, which is 'b'.  
`m[0]` a pointer to the first element of `m`, which is the array { 7, 8, 9 }  
`m[0][2]` the third element of the first element of `m`, which is 9  
`m[0][4]` *probably* the fifth element of the first element of `m`, which is 11 (actually part of the second element of `m`)

### pointer operators (address and dereference)

A pointer represents the location in memory at which a value is stored. A pointer is also a value stored in memory: it is a special kind of unsigned integer, often the same size as a long.

There are two operators used with pointers: the *address* operator `&`, and the *dereference* operator `*`.

If `x` is a variable, `&x` is “the address at which `x` is stored.”

If `p` is a pointer variable, then `*p` is “the value stored at the address pointed to by `p`, interpreted according to the base type of `p`.” **Warning:** if `p` is NULL or points to other memory not controlled by the program, attempting to evaluate `*p` will cause the program to crash!

In the example below assume that on this machine, int are four bytes (32 bits) wide and pointers are eight bytes (64 bits) wide.

```
int x = 5;  x is declared as an int and initialized to hold
            the value 5. Assume that x is stored in the four
            bytes starting at memory location 100.

int *p;     p is declared as a pointer to int. Assume that
            p is stored in the eight bytes starting at memory
            location 104.

int **pp;   pp is a pointer to a pointer to int. Assume that
            pp is stored in the eight bytes starting at
            memory location 112.

p = &x;     p is assigned the value 100, the address of x
pp = &p;     pp is assigned the value 104, the address of p
```

After this initialization, the expressions below will have the values and types indicated:

<code>x</code>	5 (int)
<code>p</code>	100 (pointer to int or int *), the address of x
<code>*p</code>	5 (int) the value at the location pointed to by p, interpreted as an int
<code>pp</code>	104 (pointer to pointer to int or int **), the address of p
<code>*pp</code>	100 (pointer to int or int *), the value at the location pointed to by pp, interpreted as a pointer to int

`**pp` 5 (int), the value, interpreted as an int, stored at the value (interpreted as a pointer) that is pointed to by pp. Whew!

### pointer arithmetic

It is legal to add an integer to a pointer and to subtract an integer from a pointer. The integer is multiplied by the width of the pointer's base type before being added. For example, if `p` is a pointer to int with the value 100, and the int type is four bytes wide, then `p+1` has the value 104 and `p-3` has the value 88.

### function call expression

The value of a called function is its returned value, which will match the function's declared return type (possibly void). When a function is called, the caller must provide actual values that are type-compatible with each of the function's declared parameters.

*expr* ( *actual1*, *actual2*, ... )

*expr* will normally be a function name, but can be any expression that evaluates to a function. Evaluates *actual1*, *actual2*, etc, in order, then causes the function to be executed, substituting the values of the actuals for the function's parameters.

The C standard library provides a large number of built-in functions; see the *C Reference Card* for a partial list.

### ternary conditional expression ( *a ? b : c* )

*condition* ? *expression1* : *expression2*

Evaluates *condition*. If *condition* is true, evaluates *expression1*, which becomes the expression's value. Otherwise, evaluates *expression2*, and that becomes the value.

For example, assuming `x`, `y` and `z` are all declared as int...

```
x = y > z ? y : z ;
```

is equivalent to

```
if (y > z) {
    x = y;
} else {
    x = z;
}
```

---

## Statements

### expression statements

Any expression can be made into a statement by appending a semicolon. The course coding standard requires one expression statement per line; however, the compiler will allow multiple statements on the same line.

All of the following are legal C statements, assuming the variables and functions shown are appropriately declared. However, the last statement has no effect and will be flagged with a warning by most compilers.

```
a += 3;
printf("%d\n", a);
a * 97;
```

## block statement (sequence)

An opening curly brace, a series of statements, and a closing curly brace is considered a single statement. The contained statements are executed in order.

```
{
    statement
    statement statement
}
```

## if and if-else statements (selection)

Allow different statements to be executed depending on the truth value of a given *condition*.

```
if ( condition )
    statement
```

If *condition* is true, execute *statement*. Otherwise, skip *statement*.

```
if (condition)
    true-statement
else
    false-statement
```

If *condition* is true, execute *true-statement*. Otherwise, execute *false-statement*.

The course coding standard requires that the contained *statement(s)* in an if are **block statements** with opening and closing curly braces. For example (assume all variables are declared)

```
if (balance >= request) {
    balance = balance - request;
    printf("Here is your money\n");
} else {
    printf("You don't have that much\n");
}
```

## switch statement (selection)

A generalized if. In an if, the condition expression is interpreted as a truth value and there are only two possible branches: true and false. In a switch, the expression may be any **integer type** and there can be one branch for each possible value of the type. The general form of switch is:

```
switch ( expression ) {
    case label1 :
        series of statements
        break;
    case label2 :
        series of statements
        break;
    (other cases as needed)
    default :
        series of statements
}
```

The *expression* must evaluate to an integer type. Each of the case *labels* must be a constant integer different from all other *labels*. Inside each case we can have a *series of statements* — which is **weird**, because in all other C control structures the body must consist of a single statement (which can be a block statement using curly braces). Each case will normally end with a break.

When execution arrives at a switch, the *expression* is evaluated. Then each *label* is compared against the *expression's*

value. If there is a match, the *statements* inside the matching case are executed; if there is no match and the optional default case is present, the default *statements* are executed.

If the *statements* end with **break**, control transfers to the point just after the switch statement's closing curly brace.

If the *statements* do not end with a **break**, control “falls through” to the *statements* in the next case, and execution proceeds until either a **break** is encountered or the switch ends. **Leaving out break is almost always an error!**

Here is an example. Assume input has type char.

```
switch (input) {
    case 'a':
        printf("handle a ");
    case 'b':
        printf("handle b ");
        break;
    default:
        printf("handle all others\n");
}
```

If input has the value 'a', this will display “handle a handle b”, since case 'a' contains no break.

## while statement (loop)

Allows a *statement* (in this course, a block statement) to be executed repeatedly, as long as a *condition* remains true.

```
while (condition)
    statement
```

If *condition* is false, skip *statement*. If *condition* is true, execute *statement* once, then re-evaluate *condition* and proceed as indicated above, repeatedly executing *statement* then re-evaluating *condition* until *condition* becomes false.

## do-while statement (loop)

Operates just like the while statement, except that the *condition* is checked after the first execution of contained *statement*, not before.

```
do
    statement
while (condition);
```

Execute *statement*. Evaluate *condition*, and if it is true, execute *statement* again and re-evaluate *condition*. Repeat this process until *condition* evaluates to false. Note the closing semicolon after the *condition*.

## for statement (loop)

Allows a *statement* (in this course, a block statement) to be executed repeatedly as long as a *condition* remains true. Provides an *initialization* and a post-statement *action*.

```
for ( initialization ; condition ; action )
    statement
```

Execute *initialization*. Evaluate *condition*, and if it is true, execute *statement*, execute *action*, and re-evaluate *condition*. Repeat this process until *condition* evaluates to false.

Example:

```
int sum = 0;
int prod = 1;
for (int val = 1; val < 10; val += 2) {
    sum += val;
```

```
    prod *= val;
}
```

After executing this code, `sum` will have the value 25, which is the sum of the numbers 1, 3, 5, 7 and 9, and `prod` will have the value 945, the product of those numbers.

## break

Executing a `break` statement causes control to exit from the nearest enclosing loop (`while`, `do-while` or `for`) and proceed to the next statement following the loop. It is also used with a slightly different meaning in the `switch` statement.

## continue

Executing a `continue` statement causes execution to skip over any remaining statements in the nearest enclosing loop's body, exactly as if control had passed to the point immediately prior to the loop body's closing curly brace. In `while` and `do-while` loops this causes the loop condition to be re-tested. In `for` loops this causes the loop action to be executed, then the condition to be tested.

---

## Functions

A *function* is a computational element that takes zero or more *parameters* and optionally returns a value. Function types are indicated using parentheses ( ).

### function prototype (declaration)

All functions in C programs must be *declared* before use. A function declaration, called a prototype, states that a function with a given name, parameter types and return type exists, but does not define the function. A **function definition** implicitly declares a function, but it is normally preferable to separate declaration from definition. A function prototype has the form:

```
return-type function-name ( parameter-list );
```

The *return-type* may be any C type. If the function returns no value, the return type is `void`. The *function-name* must be a legal identifier. The *parameter-list* is a comma-separated list of parameters, each of which has the form of a variable declaration. Here are some examples:

```
int next(void);
    next is a function that takes no parameters
    (void) and returns an int

int multiply(int x, int y);
    multiply is a function that takes two int
    parameters called x and y and returns an int

int multiply(int, int);
    a legal but discouraged equivalent to the
    previous declaration; don't do this

void output(char *s);
    output is a function that takes a pointer to a
    char (i.e., a string) and doesn't return a value

void output(const char *s);
    as above, but any attempt in the body of output
    to modify the string pointed to by s will be
    flagged as illegal by the compiler; guarantees to
    callers that the function output will not modify
    the string pointed to by s
```

## function definition

A function definition provides the code that implements a function. It has the form

```
return-type function-name ( parameter-list ) {
    local variable declarations
    statements
}
```

If the function has been previously declared, the first line must match the prototype, including the number, order, and types of parameters. It is legal to use different parameter names in a function prototype and its definition, or to omit the names entirely in the prototype, but this is bad practice and not acceptable in this course.

Any variables declared in the function are *local*: visible only within the body of the function. Each time the function is called, local variables are re-allocated and re-initialized. The only exception are **static variables**, see below.

The function executes until either a `return` statement is encountered (see below) or the end of the function body is reached.

### return statement

The `return` statement causes control to return from a function to the place from which it was called. There may be multiple `return` statements in a function.

If the function's declared return type is `void`, then the statement must have the form

```
return;
```

It is not necessary to provide a `return` statement at the end of a `void` function.

If the function's return type is anything other than `void` then the form

```
return value ;
```

must be used, and *value* must be an expression whose type is compatible with the declared return type.

### static variables

The prefix `static` may be used to declare that a local variable in a function retains its value for the duration of the program, rather than the duration of the function call. Static variables are initialized at compile time. The following function will return 101 the first time it is called, 102 the second time, 103 the third time, and so on.

```
int next_serial(void) {
    static int counter = 100;
    counter++;
    return counter;
}
```

---



---

## Comments

Comments are text embedded in your code intended for human readers rather than the compiler. C provides single-line comments and block comments.

Where two slashes together (//) appear in your C code anywhere outside a character string, the slashes and the remainder of the line on which they appear are a comment.

```
// This is a comment. The last line of this
// example has code followed by a comment.
aqlue = 42; // The Adams number.
```

Block comments start with the characters /\* and end with the characters \*/. Everything between is a comment.

```
/* This is a comment
which
continues down
here */
```

---

## Memory management

### sizeof

The sizeof operator can be applied to any type or declared variable and will evaluate to the number of bytes required to store an object of that type. sizeof operates at *compile time*, so it cannot determine the size of dynamically dimensioned objects. In the following, assume an int occupies four bytes.

```
sizeof(int) evaluates to 4
int x;
sizeof(x)    also 4
int y[5];
sizeof(y)    evaluates to 20 (five elements by four bytes)
              since the size of y was specified at compile time
```

### malloc

The malloc (“memory allocate”) function is used to request memory dynamically at runtime. malloc takes one parameter, which is the size of memory to request, in bytes. It returns a void pointer to the memory allocated, or NULL if the allocation failed, e.g. there is no more memory available. The allocated memory is not zeroed or otherwise initialized.

The following example requests memory suitable for the Client type from the typedef definition and use section.

```
Client *new_client;
new_client = malloc(sizeof(Client));
if (new_client != NULL) {
    success, initialize fields
} else {
    failure, handle out of memory error
}
```

### calloc

The calloc (“contiguous allocate”) function is similar to malloc but used where space is required for an array. It takes the number of elements and the size of a single element as parameters and returns a void pointer. For example, the following requests space for an array of 40 int values.

```
int *new_array = calloc(40, sizeof(int));
```

Like malloc, calloc will return NULL on failure. Unlike malloc, if calloc succeeds it does set all bits in the allocated memory to zero.

### free

Memory allocated by malloc and calloc must be explicitly released using free when it is no longer required. Free takes a pointer to previously-allocated memory as a parameter. Failure to free unneeded memory is called a *memory leak* and can result in performance problems and program crashes.

For example, to release the memory allocated with calloc above, we would call free(new\_array)

---

## printf and sprintf

The printf and sprintf functions (found in stdio.h) combine a format specification string with (optional) values to produce a result string. printf displays its result on the console; sprintf puts its result in a character array buffer, which you must provide.

```
printf(format, v1, v2, ...)
sprintf(buffer, format, v1, v2, ...)
```

For example,

```
int n = 2;
double bal = 291.57923;
printf("In %d accounts\nyou have $%.2f\n", n, bal);
```

will display

```
In 2 accounts
you have $291.58
```

The format string can include three kinds of things: literal characters, interpreted character combinations (also known as “escaped characters”), and value placeholders.

### literal characters

Literal characters are passed through to the result unchanged. In the example’s format string the literal characters are “In ”, “ accounts” and “you have \$”.

### escape sequences

See the section on **escape sequence character literals**, which have special meanings in format strings. Here, the only escape sequence is \n (used twice) meaning “go to the next line.”

### value placeholders (conversion specifiers)

Value placeholders (also known as conversion specifiers) specify where the provided values v1, v2, etc. are to appear in the result string as well as how they should be displayed. Values are inserted into placeholders in left-to-right order and there must be exactly one placeholder for each value.

In the example, the placeholders are %d and %.2f, which match the provided values n and bal. The %d means “display as an integer.” Similarly, %f means “display as a floating point number”. The 0.2 in %.2f means “as wide as necessary, with two digits after the decimal.”

For a summary of the available C conversion specifiers, see the *C Reference Card* section on Input/Output, particularly Codes for Formatted I/O.

---

## Copyright and use

© Greg Phillips, [greg.phillips@rmc.ca](mailto:greg.phillips@rmc.ca). This guide may be used for any purpose within the Government of Canada. All other use is governed by a Creative Commons Attribution-ShareAlike 4.0 International License. For details see <http://creativecommons.org/licenses/by-sa/4.0/>